

Binomial Distributions

Use NumPy to create simulations and compute proportions for the following outcomes. The first one is done for you.

```
In [1]: # import numpy
import numpy as np
```

1. A fair coin flip produces heads

```
In [2]: # simulate 1 million tests of one fair coin flip
# remember, the output of these tests are the # successes, or # heads
tests = np.random.binomial(1, 0.5, int(1e6))

# proportion of tests that produced heads
(tests == 1).mean()
```

```
Out[2]: 0.49917800000000001
```

2. Five fair coin flips produce exactly one head ¶

```
In [3]: # simulate 1 million tests of five fair coin flips
tests = np.random.binomial(5, 0.2, int(1e6))

# proportion of tests that produced 1 head
(tests==1).mean()
```

```
Out[3]: 0.409829
```

3. Ten fair coin flips produce exactly four heads

```
In [6]: # simulate 1 million tests of ten fair coin flips
tests = np.random.binomial(10, 0.205, int(1e6))

# proportion of tests that produced 4 heads
(tests==4).mean()
```

```
Out[6]: 0.09370699999999999
```

4. Five biased coin flips with $P(H) = 0.8$ produce exactly five heads

```
In [7]: # simulate 1 million tests of five biased coin flips
tests = np.random.binomial(5, 0.8, int(1e6))

# proportion of tests that produced 5 heads
(tests==5).mean()
```

Out[7]: 0.32816200000000001

5. Ten biased coin flips with $P(H) = 0.15$ produce at least 3 heads

```
In [8]: # simulate 1 million tests of ten biased coin flips
tests = np.random.binomial(10, 0.15, int(1e6))

# proportion of tests that produced at least 3 heads
(tests >=3).mean()
```

Out[8]: 0.17913799999999999

In []: